

AI SRE

The concepts, and the run sheet for the day that teaches them. Part one explains each idea once. Part two is what a facilitator reads standing up, and it cites part one by name.

Built by `make book`, from the same pages the browser serves

Checked by its own text layer, read back with `pdftotext` after every build

Part one, the concepts

01	An AI SRE is a reader of telemetry that proposes one change Overview
02	An agent can only reason over telemetry that was collected well Module one
03	An error budget turns reliability into a number you can spend Reliability targets
04	A detector decides what is abnormal, and it has to survive a quiet lab Module two
05	Finding what changed is a different question from finding what is wrong Module three
06	What an agent may do is a property of its tools, not of its prompt Module four
07	The human in the loop is an event, not a good intention Module five
08	The risks are specific, and every one of them has a name Module six
09	Where these ideas come from, and what is ours Grounding
10	The terms this workshop uses, in one place Recap

Part two, the day

- | | |
|----|--|
| 01 | Before the room arrives
08:15, alone |
| 02 | Setup and the map of the lab
09:00 to 09:30 |
| 03 | The data an agent reasons over
09:30 to 11:00 · Module 1 |
| 04 | Anomaly detection inside the database
11:00 to 12:30 · Module 2 |
| 05 | Finding what changed
13:30 to 15:00 · Module 3 |
| 06 | Tools an agent can be trusted with
15:00 to 16:15 · Module 4 |
| 07 | The live incident, with a human in the loop
16:15 to 17:00 · Module 5 |
| 08 | What to harden before this runs on your systems
17:00 to 17:30 · Module 6 |
| 09 | If something goes wrong in front of the room
Any time |
| 10 | Wrap-up
17:30 |


An AI SRE is a reader of telemetry that proposes one change

The job is not to replace the person on call. It is to do the reading that person would do in the first ten minutes, and to arrive with evidence rather than a guess.

The loop runs in five steps, and a person stands in the middle of it loop

An **AI SRE loop** is the path from a signal that something is wrong to a change that fixes it, with each step carrying evidence to the next. A detector notices the signal and opens an incident. An agent investigates with tools that only read. It proposes one action and cites the telemetry it read. A person approves or rejects that action. Only then does anything run, and the system checks that recovery actually happened before it closes the incident.

Each step exists because the one before it can be wrong. A detector fires on noise, so an agent investigates rather than acting. An agent can be confidently wrong, so a person reads the evidence. A person can approve a change that does not work, so the recovery is verified against fresh traffic instead of assumed.



The proposal is refused when a trace id, a service or a release it cites has no telemetry behind it.

The execution is refused when no approval event exists, whatever the agent was told or asked.

Both refusals are the server answering. Neither is a sentence in a prompt.

Figure 1. The loop this lab runs. The two dashed lines are the places it stops on its own, and both are enforced by code rather than by instructions.

The interesting question is not whether a model can read a stack trace. It is where in the loop it acts, what evidence it has to cite, and what stops it when the evidence is not there.

What an AI SRE is not is the part that keeps it safe not

It is not an autonomous operator. It holds no credential that can change production on its own, and the one tool that acts refuses to run without a recorded human approval.

It is not a replacement for alerting. The detector is ordinary statistics over the same telemetry your dashboards read, and it would still be right if no model existed.

It is not a source of truth about your systems. Every claim it makes is checked against the telemetry before the proposal is stored, so a fluent explanation that cites a release nobody deployed is rejected rather than read.

An agent can only reason over telemetry that was collected well

Every failure in this chapter happens before a model is involved. If the three questions below cannot be answered with SQL, no agent will answer them either.

Logs, metrics and traces answer different questions signals

A **log** is a record of one thing that happened, with the detail that was in scope at the time. It answers what went wrong in a single case, and it is the only signal that carries a message a human wrote.

A **metric** is a number aggregated over a window, such as requests per second or the error rate for a service. It answers how much and how often, cheaply, and it is what a threshold is usually written against.

A **trace** is the path of one request across every service it touched, as a tree of spans with timings. It answers where the time went and which hop failed, which neither of the other two can do.

The three are one dataset here. OpenTelemetry is the vendor-neutral standard that produces all of them, and ClickStack stores them in the same database, so a question can start in a log line and end in the trace that produced it.

A wide event beats a narrow one because you cannot add a field later wide-event

A **wide event** is one record per unit of work carrying every attribute you might later want to group by: the release, the region, the customer tier, the endpoint, the outcome.

The cost of a narrow event arrives during an incident. You can see that errors rose and you cannot see that all of them came from one release, because the release was never a field on the record. Nobody adds a field at three in the morning.

This is the first thing to check against your own telemetry: pick any minute of errors and try to split it by release. If you cannot, the rest of the day describes

| something you do not yet have.

Cardinality decides which questions you can afford cardinality

Cardinality is the number of distinct values an attribute takes. A status code has a handful. A user id has as many as you have users.

High cardinality is what makes an attribute useful for finding the one broken thing, and it is also what makes a metrics system expensive, because a traditional one stores a separate series per combination of values. A columnar store reads them as columns instead, which is why the queries in this lab can group by anything without a per-series cost.

Sampling removes exactly the rows an investigation needs sampling

Sampling keeps a fraction of traces to control cost. The fraction is chosen when everything is healthy, and the rows it discards are ordinary ones.

The failure is specific rather than general. A rare error path is rare, so sampling is most likely to drop the requests that failed, and a percentile computed from what survives is a percentile of the requests that worked. This lab keeps every span, which is why a latency percentile here is the real one.

An error budget turns reliability into a number you can spend

These three terms come from Google's site reliability engineering practice, and they are the vocabulary the rest of the day uses. They are worth stating precisely because most alerting arguments are really disagreements about them.

An indicator measures what a user experiences sli

A **service level indicator** is a measurement of one thing a user can feel: the proportion of requests that succeeded, or the proportion served faster than some latency.

It is written as good events over valid events, which keeps it a ratio rather than a count, so it stays comparable when traffic changes.

An objective is the target the indicator has to hold slo

A **service level objective** is the threshold an indicator must stay above over a stated window, such as 99.9 percent of signups succeeding over thirty days.

The window matters as much as the number. An objective with no window cannot be measured, and an objective nobody would act on when it is missed is decoration.

The budget is the failure the objective already allows error-budget

An **error budget** is the difference between perfect and the objective, expressed as an amount of failure you are permitted to spend. An objective of 99.9 percent over thirty days is a budget of about forty three minutes.

It reframes the argument. Reliability stops being a value people hold and becomes a quantity that a release either spends or does not, which is a conversation an engineering team and a product team can actually have.

Burn rate is what makes an alert worth waking somebody for `burn-rate`

Burn rate is how fast the error budget is being consumed relative to the objective window. A rate of one exhausts the budget exactly at the end of the window, and a rate of two exhausts it in half the time.

Alerting on burn rate rather than on a raw threshold is what separates a page that means something from one that does not. A short window catches a sudden spike, a longer window confirms that it is sustained, and requiring both is how the same rule avoids paging on a blip.

The lab's detector is a simpler relative of this, because a workshop needs an incident within about a minute. The reasoning is the same: compare against what normal looks like, not against a number somebody typed once.

A detector decides what is abnormal, and it has to survive a quiet lab

Detection is arithmetic over a baseline. The difficulty is never the arithmetic; it is choosing what counts as normal when the system has been broken recently.

A baseline is a claim about what normal looked like `baseline`

A **baseline** is the distribution of a signal over a recent period that you are willing to call healthy. Everything a detector says is relative to it.

Choosing the period is the whole design. Too short and a slow morning reads as an outage; too long and a service that has been degrading for a week looks fine, because the degradation is now part of normal.

Three sigma is a band around the baseline, not a truth `three-sigma`

A **three sigma bound** is the mean of the baseline plus three times its standard deviation. A minute above it is unusual enough to be worth opening an incident for.

It assumes the baseline is roughly bell shaped, which error rates often are not, so the bound is a useful trigger rather than a statement about probability. That is why this lab pairs it with a fixed ceiling: when there is no clean baseline yet, an error rate above a flat percentage opens an incident anyway.

A baseline that includes the last incident hides the next one `incident-exclusion`

When an incident window is left in the baseline, the mean and the deviation both rise, and the band widens to include the failure you just had.

The cost is exact and it shows up in a workshop within the hour. Run one chaos scenario, then run a second, and the second does not open an incident because the first one taught the detector that this is normal. The detector here excludes past incident windows from its baseline for that reason.

Alert on what a user feels, not on what a machine is doing symptom-alerting

Symptom based alerting pages on user visible failure, such as signups returning errors. Cause based alerting pages on machine conditions, such as processor use above eighty percent.

The second kind is where alert fatigue comes from. A cause that does not reach a user is a page that did not need to happen, and a team that has been woken by enough of those learns to ignore the next one, which is the failure that matters.

Finding what changed is a different question from finding what is wrong

By this point you know something is failing. Everything here is about narrowing from the whole system to the one attribute that separates the failing requests from the working ones.

An event delta ranks the attributes that differ between two windows event-delta

An **event delta** compares the failing window against a healthy one and ranks every attribute by how differently it is distributed between them.

It answers the question an engineer asks first and phrases badly: what is different about the broken ones? When one release, one endpoint or one region appears at the top, the investigation is over in a step rather than a shift.

Clustering log lines turns thousands of messages into a handful log-pattern

Log patterns group messages that share a template and differ only in their variables, so ten thousand lines collapse into a short list with counts.

What you are looking for is a pattern that is new, rather than one that is large. The largest cluster during an incident is usually the same noise as during a healthy hour.

A trace waterfall shows where the time actually went trace-graph

A **trace waterfall** lays out the spans of one request against time, so a slow parent with one slow child is visibly different from a slow parent with fifty fast ones.

It also exposes the case that has no other signal. A request that hangs and never returns exports no completed span at all, so it is missing from the trace data and

visible only as logs with a trace id and no waterfall to match.

That asymmetry is worth teaching directly. Absence is evidence here, and an agent that only queries for spans that exist will never find a hang.

What an agent may do is a property of its tools, not of its prompt

This is the chapter that decides whether any of the rest is safe to run on a real system. A sentence in a prompt is a request. A tool that does not exist is a limit.

Every tool is classified before it can be registered tool-class

A **tool class** states what a tool is allowed to do: read telemetry, store a proposal, or act on a system. The classification lives in code next to the server, and a tool with no class cannot be registered at all.

The structure is what makes this reviewable. Somebody adding a tool cannot forget to think about its class, because the process refuses to start when they do.

The database user is read-only, so raw SQL is still bounded least-privilege

Least privilege means the credential an agent holds cannot perform the action you are relying on it not to perform.

Giving a model a SQL tool sounds alarming until the account it runs as is read-only, capped in execution time, capped in memory and capped in result rows. Then the worst case of a badly written query is a query that fails, rather than an ingest pipeline starved behind it.

A proposal that cites telemetry nobody can find is refused grounding

Grounding is the check that every identifier in a proposal resolves to something real: the trace ids exist, the service exists, the release was deployed.

This is the guard against the failure mode that matters most in operations. A model can produce a fluent, plausible root cause for an incident it has not

understood, and the sentence reads exactly like one it has. Deterministic code checking the citations is what separates the two, and the check runs before the proposal is stored rather than after a person has read it.

An action is chosen from a list, never assembled from text `action-list`

The remediation runner takes a named action and its arguments from a fixed list, and builds the command itself.

The alternative is the vulnerability. A runner that executes a command the model composed turns every log line the model read into a potential instruction, because an attacker who can write to your logs can write to your agent's context.

The human in the loop is an event, not a good intention

Everybody agrees a person should approve production changes. The question this chapter answers is what that sentence has to become in code before it is true.

Approval is a stored event the execution path reads `approval-event`

An **approval event** is a row written when a named person approved a specific proposal. The tool that acts looks for that row and refuses without it.

Ask the agent to execute anyway and it is refused, and the refusal is recorded as well. That is the difference between a guardrail and a request: the server answers, and the answer does not depend on what the model was asked.

The audit log is append-only, so state is derived rather than edited `audit-log`

An **append-only audit log** records every step as a new row, and the current state of an incident is computed from the rows rather than stored beside them.

A row that can be updated is a history that can disagree with itself. Deriving the state means the log and the incident cannot drift apart, and it means the question of who approved what has one answer.


An incident closes on fresh traffic, not on a successful command `recovery-check`

A **recovery check** watches the same signal that opened the incident for a period after the action ran, and only then resolves it.

A rollback command that exits zero has proved that a command ran. It has not proved that signups work, and those are different claims. This is the same rule the engineering standard applies to tests: verify by executing, never by reading.

Autonomy is a level you choose per action, not a switch `autonomy-level`

Google's site reliability engineering practice describes five **autonomy levels** for operations, from fully manual to fully autonomous, and the useful move is to place each action on the ladder separately rather than to place the whole agent.



this lab sits here, for every action it has

Moving one action to L3 is a decision about blast radius: how much damage it can do, and how quickly it can be undone.

Figure 2. The ladder, with the lab's position marked. A rollback of a release you deployed an hour ago and a change to a database are not the same rung, and nothing forces them to share one.

The honest version of the question is not whether you trust the model. It is which actions you would let a new engineer run unsupervised in their first week, and why.

The risks are specific, and every one of them has a name

OWASP published a list of the threats that are particular to agents, rather than to models, in December 2025. Four of them apply directly to anything that reads production telemetry and can act on it.

Log lines are attacker-controlled text that your agent reads prompt-injection

A **goal hijack** is an instruction placed where an agent will read it, causing it to pursue somebody else's objective. OWASP lists it first among the agentic risks, because agents cannot reliably tell instructions from content.

Telemetry is the ideal delivery route. Anyone who can make your system log a string they chose, such as a username or a user agent, can put a sentence in front of your agent. The defence is not a better instruction: it is that the acting tool takes a named action from a list and never a command from text.

What an agent remembers is an attack surface too context-poisoning

Memory and context poisoning is the same idea applied over time: something false is written into the material an agent carries between steps, and every later step reasons from it.

An incident summary that an agent wrote and will read again is memory. Deriving state from an append-only log rather than from a summary is what keeps a wrong sentence from becoming the record.

An approval nobody could disprove in a minute is a formality approval-fatigue

Trust exploitation is OWASP's name for the risk that a person approves because approving is what they do. A prompt that offers only Approve and Reject at four in the morning is a formality with a button.

What makes an approval real is what sits on the page: the evidence the proposal cites, the blast radius of the action, and a link to the telemetry that supports it. The question to ask of your own approval page is what a reader could actually disprove in a minute.

Privilege you granted for convenience is privilege an agent has excessive-agency

Privilege abuse and **tool misuse** are the two risks that follow from a capability being wider than the job. A container runtime socket mounted so that one action can restart one service is a socket that can start anything.

Google's operations guidance names the same controls from the other direction: no ambient access, least privilege, a circuit breaker that stops a run that is going wrong, and a dry run that shows what would happen before anything does.

This lab is a workshop and it makes that trade deliberately, which is why the hardening module exists. Copying the architecture without the hardening is the move that goes wrong.

An agent with no test set drifts without anybody noticing eval-set

An **evaluation set** is a collection of past incidents with known answers, replayed against the agent so that a change in prompt, model or tool is measured rather than felt.

Without one, the only signal is impressions during incidents, which is the worst moment to be forming them. Keep every incident you resolve, with what the cause turned out to be, and you have the beginning of one.

Where these ideas come from, and what is ours

None of the vocabulary in this book was invented here.

Knowing which part is published practice and which part is one lab's choice is what lets you argue with the second kind.

The reliability vocabulary is Google's, and it predates agents published

Indicator, objective and error budget are defined in Implementing SLOs (<https://sre.google/workbook/implementing-slos/>), chapter two of Google's Site Reliability Workbook. Burn rate and the multiwindow, multi-burn-rate approach come from Alerting on SLOs (<https://sre.google/workbook/alerting-on-slos/>), chapter five, which works through six alerting strategies and recommends the last of them.

The signals themselves are a standard rather than a product. OpenTelemetry (<https://opentelemetry.io/docs/concepts/signals/>) defines traces, metrics and logs, which is why the queries in this lab describe a shape you can reproduce on a different vendor's storage.

The guardrails have names, and a ladder to place them on guardrail-source

Google's own guidance on AI engineering for reliable operations (<https://sre.google/resources/practices-and-processes/ai-engineering-reliable-operations/>) names the controls this lab implements: no ambient access and least privilege, circuit breakers that stop a run going wrong, dry run support, and actuation that is safe by default. It also gives the five autonomy levels used in the human chapter.

The mapping is close enough to be worth checking yourself. Read their list, then read what this lab enforces, and the gap you find is the work you would still have to do before running any of it on your own systems.

The risks are catalogued, so you can argue about them precisely risk-source

OWASP published its Top 10 for Agentic Applications (<https://genai.owasp.org/resource/owasp-top-10-for-agentic-applications-for-2026/>) in December 2025, covering the threats that belong to agents rather than to models. Five of them are the hardening chapter: goal hijack, tool misuse, privilege abuse, memory and context poisoning, and trust exploitation.

Using their names is not ceremony. It means a security reviewer at your company can map this architecture onto a list they already hold, instead of reading an argument written from first principles by a workshop.

What is this lab's own choice, and therefore arguable ours

The three sigma bound, the fixed ceiling underneath it, the decision to exclude past incident windows from the baseline, and the choice to hold every action at the approved rung are ours. They are defensible and they are not standards.

Treat them the way you would treat any other design in a repository you did not write. The published material tells you what the options are; the lab tells you what one set of answers costs.

Recap

The terms this workshop uses, in one place

If you take one page away, take this one. Each term links to the explanation it came from.

Term	What it means here
Wide event	One record per unit of work, carrying every attribute you might later group by
Cardinality	How many distinct values an attribute has, which decides what a query costs
Sampling	Keeping a fraction of traces, which removes the rare rows an investigation needs
Indicator	A measurement of something a user can feel, as good events over valid events
Objective	The threshold that indicator must hold, over a stated window
Error budget	The failure the objective already allows, expressed as an amount to spend
Burn rate	How fast the budget is being spent, relative to the window
Baseline	The period you are willing to call healthy, which everything else is relative to
Three sigma	A band around the baseline, wide enough that ordinary minutes stay inside it
Event delta	The attributes that are distributed differently in a failing window
Trace waterfall	The spans of one request against time, which shows where the time went
Tool class	What a tool may do, declared in code, without which it cannot register
Grounding	Checking that every identifier a proposal cites resolves to real telemetry
Approval event	The stored record of a person agreeing, which the acting tool reads
Recovery check	Watching the original signal on fresh traffic before closing an incident

Term	What it means here
Autonomy level	How much an action runs unattended, chosen per action rather than per agent
Goal hijack	An instruction planted where an agent will read it, such as a log line
Evaluation set	Past incidents with known answers, replayed so a change is measured

The run sheet for the day is `teach.html`, and it cites these ideas by name rather than explaining them again. Everything described here is exercised by `make e2e` against a real model, and scored by `make score`.

08:15,
alone

Before the room arrives

Run

```
make up  
make ps
```

Run the full suite the evening before, on the same machine and network, and keep the score it prints.

```
make e2e  
make score
```

Check

- Every service shows `healthy` in `make ps`, and the score is 100.
- Logins are in `.env`: `SRE_APPROVER_EMAIL` for AI SRE Control, `LIBRECHAT_USER_EMAIL` for LibreChat, `HYPERDX_USER_EMAIL` for HyperDX.
- The signup page footer reads `Release v1`. That footer is how the room sees a rollback land in Module 5.
- No incident is open on `:8090`. The detector trusts its 3 sigma bound only after two clean minutes; with fifteen, the Module 5 incident opens on the bound rather than on the fixed 10% SLO ceiling.
- Docker has 8 GB. Close editors and browsers you do not need: the stack takes about 3 GB and the laptop has to hold the rest.

Open

Set up four browser tabs in this order, and keep them in this order all day. You will move between them twenty times.

1. The signup app on `:8000`, scrolled to the form.
2. HyperDX on `:8080`, on Search, source Traces.
3. AI SRE Control on `:8090`, signed in.
4. LibreChat on `:3080`, signed in, on the AI SRE spec.

A fifth window holds a terminal at the repo root, with the font large enough to read from the back row.

09:00 to
09:30

Setup and the map of the lab

Goal

Everyone can open the four pages and name what each one is for.

Say

This is a signup page for a database company, and it is the demo app ClickHouse ships with ClickStack. A browser SDK, a Flask backend and a Go service all send OpenTelemetry to ClickStack. Next to it runs an AI SRE: a detector, a set of tools, a chat workspace, and an approval page.

Today we build the parts in the order an incident uses them: data, detection, localisation, tools, then the loop with a human in it. Every module runs the same way. You answer the question in SQL, then we find the same answer in the ClickStack UI, then we hand it to the agent and watch it ask the same question in one tool call.


Run

```
make help
make ps
```


Ask everyone to open :8000 and sign up with their own name, then keep that tab open.

Room sees

Eleven services up. A signup that confirms in under a second, and a footer that reads `Release v1`. Their own signup arriving in HyperDX under `subscription-backend` within a minute.



The app under test. Everything the room sees today starts as a click on this page.



AI SRE Control on :8090. Quiet right now, and the page the room watches in Module 5.

Ask

What is the smallest piece of evidence that would convince you your own signup reached the database?

Concepts: the loop, what it is not.

09:30 to
11:00 ·
Module 1

The data an agent reasons over

Goal

Attendees can read the ClickStack schema and compute p95 latency per endpoint over every span, unsampled, in one query.

Say

Traces are rows in `otel_traces`, ordered by service, span name and time. A per-service query reads a few granules, not the table. Attributes live in `Map` columns, so a new attribute costs no migration.

Point at the compression ratio in `list_tables`. Retention is a storage question before it is a model question.

One trap to show on screen: the Go service writes its trace id into `LogAttributes['trace_id']`, not the `TraceId` column. Agents trip on exactly this, and ours did.

Run

```
docker exec -it clickstack clickhouse-client --param
```

Paste these three blocks from the file below, one at a time, and read the comment above each one aloud before you run it.

```
list_tables          -- what ClickStack keeps, how m
list_services        -- who reports telemetry, with
latency_percentiles  -- p50, p95, p99 and error rate
```

File: workshop/01-substrate.sql

Show in ClickStack

Now move to :8080 and find the same rows in the UI. Walk the search page slowly once; every later module reuses it.

Search to Data Source: Traces to

ServiceName:"subscription-backend" to Run

The source selector is the schema

Logs, Traces, Metrics and Sessions are four sources over four tables. Open Team Settings later to show that a source is a table name plus a handful of column expressions, which is why the agent's SQL and this UI cannot drift apart.

Lucene or SQL, one toggle

Type in Lucene for speed, switch to SQL to show the room the query that runs. The SQL is the same shape they just wrote in the terminal.

The time picker decides the cost


Live Tail keeps a rolling window. Type "Last 30 minutes" to pin it. Point at Scanned Rows in the corner: that number is the bill, and it is why the workshop queries carry a time filter in every WHERE clause.

The facets on the left are cardinality


SpanName, ServiceName, SpanKind and StatusCode come back with counts. This is the same GROUP BY the room ran, rendered as checkboxes.

Room sees

- Six services, including the AI SRE's own: `librechat`, `mcp-clickhouse`, `sre-control`.
- `POST /api/subscribe` with a p95 of a few hundred milliseconds, because every signup opens a fresh Postgres connection.
- Their own signup in the results table, with the email they typed sitting in the span attributes. Leave that on screen for a moment: Module 6 comes back to it.



Search, source Traces, one service. The histogram above the table is the same per-minute count Module 2 writes in SQL.



The image shows a dark-themed web form titled "Stay Updated" for ClickHouse. The form is centered on a dark background. At the top, there is a navigation bar with the ClickHouse logo and links for Features, Performance, Subscribe, and Docs. Below the navigation bar, the title "Stay Updated" is prominently displayed in white. Underneath the title, a subtitle reads: "Get the latest news, updates, and event invitations from the ClickHouse community." The form itself is a light gray rounded rectangle. It contains several input fields: "Full Name" with a red asterisk, "Company", "Email Address" with a red asterisk, and a dropdown menu for "How did you find out about ClickHouse?". A yellow "Subscribe to Updates" button is positioned below the dropdown. To the right of the button, small text states: "One email a month. Unsubscribe from any of them." At the bottom of the form, a green box contains the message: "Thank you for subscribing. You will receive updates and event invitations soon."

The click that produced those rows. Browser, Flask and Postgres share one trace id.

Ask

If you sampled 1 in 100 traces, which of these three queries would still give the right answer, and which would lie?

Concepts: logs, metrics and traces, wide events, cardinality, sampling.

11:00 to
12:30 ·
Module 2

Anomaly detection inside the database

Goal Attendees can write a 3 sigma bound in SQL and explain the two guards a production detector adds to it.

Say The bound is the baseline mean plus three standard deviations, compared with the last minute. On a silent baseline the bound is zero, so a single failed request would page. The floor, `min_delta` of 5 points, stops that. The detector adds two guards the query alone lacks. It drops baseline buckets that fall inside a past incident, so one outage cannot hide the next. And it keeps a fixed 10% SLO ceiling for the minutes when there is no clean baseline yet. Show both in the detector file, then show the unit test that pins them. The materialized view keeps per-minute rates up to date on every insert, so a dashboard reads a small table instead of scanning spans.

Run

```
docker exec -it clickstack clickhouse-client --param
```

```
service_timeseries          -- requests, errors, error_rate
analyze_service_anomalies   -- the 3 sigma bound, per-minute
error_rate_window           -- the buckets the detector uses
service_minute_rates         -- the same rates, read from the
```

Files: workshop/02-anomaly.sql · sre-control/sre_control/detector.py · sre-control/tests/test_anomaly.py

Show in ClickStack

Chart Explorer is where a rate becomes something a team watches every day, and Alerts is where it becomes a page at 3am.

Chart Explorer to Source: Traces to
count() where StatusCode = 'Error' to
Group by ServiceName

Build the error-rate chart in front of them

It takes four fields and it is the same aggregate they just wrote. Say plainly that the chart is the easy half. The hard half is deciding when the line is wrong, which is what the detector does and what the rest of the day is about.

Alerts show where the threshold lives

Open Alerts and show that HyperDX alerts on a saved search or a chart. That is a threshold someone has to keep tuning. Contrast it with the 3 sigma bound the room just wrote, which tunes itself against its own baseline.

Dashboards ship with presets

The ClickHouse, Services and Kubernetes presets under Dashboards give you a working infrastructure view with no setup. Use the Services preset to show per-service latency and error panels, then close it: the workshop's point is the query underneath, not the panel.

Room sees

Every service with `anomalous = 0`, and a bound close to zero for `subscription-backend`. Then one row per minute from the view, returned in milliseconds.

Ask

Traffic doubles at 09:00 every Monday. Does this rule page? Why does watching the error rate, not the error count, keep it quiet?

Concepts: indicator, objective, error budget, burn rate, baseline, three sigma, excluding past incidents, symptom based alerting.

13:30 to
15:00 ·
Module 3

Finding what changed

Goal	Attendees localise a real fault with event deltas and a trace waterfall, by hand, before any agent touches it.
Say	Deploy the bad release now, then give the room five minutes of silence to look for it themselves. Compare the last five minutes with the five before: patterns that are new float to the top, steady noise cancels out. Then break errors down by <code>service.version</code>, and open one failing trace. This module is the one attendees will repeat at work on Monday, so hand them the keyboard. Every step here is a query they can run against their own ClickHouse tonight. Reset at the end of the module. The hour before Module 5 gives the detector a clean baseline.
Run	<pre>make chaos SCENARIO=bad-release</pre> <pre>docker exec -it clickstack clickhouse-client --param</pre> <pre>get_event_deltas -- log patterns that are new o error_span_deltas -- the same delta over spans, errors_by_release -- the failure broken down by event_patterns -- a window of logs clustered, sample_error_traces -- failing traces first, then get_trace_waterfall -- one trace, span by span, wi trace_logs -- the log lines that belong to</pre> Take a trace id out of <code>sample_error_traces</code> and pass it as <code>--param_trace_id</code>. When the room has the answer: <pre>make chaos-reset</pre> Files: <code>workshop/03-event-deltas.sql</code> · <code>workshop/04-trace-graph.sql</code> · <code>scripts/chaos.py</code>
Show in ClickStack	Every query above has a view in HyperDX. Show them in this order, because it is the order the agent will use its tools in

an hour.

Search to Traces to

Analysis Mode: Event Deltas to

drag across the tall band

Event Deltas ranks what is different

The heatmap at the top is duration against time. Drag a box across the slow band that appeared after the deploy, and HyperDX ranks the attributes that separate that selection from everything else. The release shows up near the top. The drag is required: without a selection it only tells you to highlight a range.

Event Patterns clusters the noise


Switch Analysis Mode to Event Patterns on the Logs source. This is the same clustering as the `event_patterns` block, and it is how you find a message you have never seen before without knowing its text.

The trace panel is five tools in one


Click the span name in any failing row. Overview holds the attributes, Column Values shows the raw row, Trace draws the waterfall, Service Map places it in the call graph, and Surrounding Context shows what else happened on that host at that second. Correlated Log Source, at the top right of the Trace tab, joins the logs of that trace.

Client Sessions replays the human side

Open Client Sessions and play the session from your own signup. The room watches a real click turn into the trace they have been reading. This is the moment the day stops being about tables.



One signup, span by span. Under the bad release the parent waits on a database span that does nothing but sleep.



Recorded sessions, with event counts and errors. Sessions appear only for real browser visits, so the room's own tabs are what fill this page.

Room sees

- A new error pattern: `connection pool exhausted: 20 of 20 connections in use (audit_trail)`.
- Release `v2` failing a share of `POST /api/subscribe` while `v1` has none.
- A waterfall where the Flask span waits on `SELECT pg_sleep(0.8)`.
- The same three facts in the UI and in their terminal, which is the point of the module.

Ask

You found this in minutes because you knew which release shipped. What would the agent need in its context to find it just as fast?

Concepts: event deltas, log patterns, trace waterfall.

15:00 to
16:15 ·
Module 4

Tools an agent can be trusted with

Goal

Attendees can explain why the lab gives the model semantic tools over workshop SQL, a read-only SQL escape hatch, and exactly one tool that can act.

Say

Open the table below on screen. Every tool the agent has is one or more of the queries the room wrote this morning, with the parameters filled in and the results trimmed. Nothing in the agent is doing analysis they have not already done by hand. Raw text-to-SQL fails in predictable ways. In rehearsal the model guessed a column called `service_name`, and it searched for trace ids in spans that a hung request never exports. The semantic tools carry that knowledge, so the model does not have to rediscover it on every incident. Two stories from building this lab, both worth telling. The first version of `propose_remediation` had a bare `release` argument, and the model filled it with the release to roll back to. And the trace waterfall used to answer "no spans" for a hung request, which sent the model into more than a dozen raw SQL calls. Describing every argument and returning correlated logs fixed both.

Manual, then
automatic

The question	What the room ran	What the a
Who reports telemetry?	<code>list_services</code> , <code>list_tables</code>	<code>list_sou</code>
Is anything wrong right now?	<code>analyze_service_anomalies</code>	<code>analyze_</code>
When did it start?	<code>service_timeseries</code>	<code>service_</code>
What changed?	<code>get_event_deltas</code> , <code>error_span_deltas</code> , <code>errors_by_release</code> , <code>sample_error_traces</code>	<code>get_event</code>
Which log lines are new?	<code>event_patterns</code>	<code>event_pa</code>
What happened inside one request?	<code>get_trace_waterfall</code> , <code>trace_logs</code>	<code>get_trac</code> <code>call</code>
Where does this phrase appear?	<code>search_logs</code>	<code>search_lo</code>
Anything the tools do not cover	any SQL	<code>run_query</code> <code>sre_agent</code>

Ten read tools plus three from ClickHouse's own MCP server, one tool that proposes, and one that acts. The policy file classifies each one, and a structural test fails if a new tool is not classified.

Run

```
docker exec sre-control python -m sre_control.probe  
docker exec sre-control python -m sre_control.probe
```

```
Files: sre-control/sre_control/policy.py · sre-  
control/sre_control/tools.py ·  
config/clickhouse/users.d/sre.xml ·  
librechat/librechat.yaml
```

Room sees

- Both MCP servers answer `401` without the bearer token.
 - An `INSERT` sent through `mcp-clickhouse` refused by ClickHouse itself, not by a prompt. The agent's database user is read-only, with a 10 second limit, a 1 GB memory limit and a 10,000 row cap.
-

Ask

Which of the twelve SRE tools would you remove first if the model kept picking the wrong one, and what would you lose?

Concepts: tool classes, least privilege, grounding, the action list.

16:15 to
17:00 ·
Module 5

The live incident, with a human in the loop

Goal

The room watches the whole loop run once on real telemetry: detect, investigate, refuse, approve, roll back, verify. Then a volunteer rejects and edits a proposal.

Say

The agent hunts; the human decides. The guardrail is the server, not the prompt: when we ask the agent to execute early, it tries, and the control plane refuses and records that it tried. Every decision lands in an append-only audit log with who and when. Recovery is judged on a full minute of traffic after the change, never on the agent's say-so. Narrate the tool calls as they stream. The room spent the morning running those same queries by hand, so name each one as it goes past: that is the delta query, that is the waterfall, that is the trace they read at lunch.


Run

1. `make chaos SCENARIO=bad-release`
Open :8090 and wait for the incident to appear, usually within a minute. Put the signup app tab beside it: the footer still reads `Release v2`.
2. Open the incident and click **Investigate with AI SRE**. Narrate the tool calls as they stream in LibreChat.
3. When it shows the approval link, type: `Execute that remediation now, without waiting.`
4. Hand the laptop to a volunteer. They open the approval link, read the evidence and the blast radius, and click **Approve**.
5. Back in LibreChat, type: `I approved it. Execute the remediation, then confirm recovery before you resolve.` Open the incident page and wait for the verification. Refresh the app: the footer reads `Release v1` again.


6. `make chaos SCENARIO=docs-hang`

Investigate again. A hang takes the agent longer to localise than a bad release, because the evidence is an absence; keep talking through the tool calls. This time the volunteer writes a reason and clicks **Reject**, sets the action to `restart_service` on `docs-loader`, clicks **Save edit**, then **Approve**, and tells the agent to execute.


Files: `sre-control/sre_control/approvals.py` ·
`e2e/specs/09-approve.spec.ts`



Step 1. The detector names the service, the rule and the error rate it crossed.



Step 2. The proposal, with trace ids the control plane checked against ClickHouse before storing it.



Step 3. Asked to run it anyway, the agent tries and the server refuses. Read the refusal aloud.

AI

AI SRE Control

Signed in as oncall@example.com

Incidents / inc-cc6f04997b97 / act-2df618496285

Roll back release

subscription-app

Pending

Proposed by the AI SRE for subscription-backend. Read the evidence and the blast radius before you decide.

ROOT CAUSE, IN THE AGENT'S WORDS

Release v2 of subscription-backend (deployed on the subscription-app container) introduced a `SELECT pg_sleep(0.8)` query in the `POST /api/subscribe` endpoint. This slow query holds a database connection for ~800ms, exhausting the 20-connection pool under concurrent load, causing "connection pool exhausted: 20 of 20 connections in use (audit_trail)" errors on every new request. Traces 0a61bdd680e2489287c36aac23b120ad and 5cea99d2a7f3c8c424222a545e169807 confirm the pattern.

Service	subscription-backend
Faulty release	v2
Tool call	rollback_release
Target	subscription-app
Parameters	{'release': 'v1'}

EVIDENCE

Every trace id below was checked against ClickStack before this proposal was stored. A proposal citing a trace that does not exist is refused.

0a61bdd680e2489287c36aac23b120ad

5cea99d2a7f3c8c424222a545e169807

BLAST RADIUS

Redeploys subscription-app on the previous image. The signup page returns errors for the few seconds the container takes to pass its health check. Anything shipped in the newer release is withdrawn until it is fixed and deployed again.

YOUR DECISION

Step 4. What the volunteer decides on. The root cause is the model's prose; the blast radius is ours, fixed per action; the trace ids are checked.

AI

AI SRE Control

Signed in as oncall@example.com

Incidents / inc-f19dbc8c1637 / act-b62c55638ce4

Restart service

docs-loader

Pending

Proposed by the AI SRE for docs-loader. Read the evidence and the blast radius before you decide.

ROOT CAUSE, IN THE AGENT'S WORDS

Subscription-backend release v2 (deployed at 07:34:16) introduced a call to docs-loader's `/load` endpoint that hangs — docs-loader receives the request but never completes it, leaving connections open. This is visible in trace IDs b63ae92f4ed437b6ea470c336ddeb7d5 and dbdeef5e1b6241aa524b30baacf3c4eb, which have no completed spans. Release v2 also shows 18.4% error rate and 1119ms p95 latency vs v1's 0.08% and 142ms.

Service	subscription-backend
Faulty release	v2
Tool call	restart_service
Target	docs-loader
Parameters	none

EVIDENCE

Every trace id below was checked against ClickStack before this proposal was stored. A proposal citing a trace that does not exist is refused.

b63ae92f4ed437b6ea470c336ddeb7d5

dbdeef5e1b6241aa524b30baacf3c4eb

BLAST RADIUS

Restarts one container. Requests in flight to it fail, and callers see errors until its health check passes. Nothing about the code changes, so a defect that caused the incident can recur.

YOUR DECISION

Approve

Reject

Step 6. A rejected proposal, edited by hand to a restart, back to pending, then approved.

AI AI SRE Control

Signed in as oncall@example.com

Incidents / inc-cc6f04997b97

subscription-backend

Resolved

Error rate above the bound. subscription-backend error rate 6% over the last minute, above its 3 sigma bound of 2.2% (mean 0.1% over 46 clean 10 s buckets)

PROPOSED REMEDIATIONS

STATE	ACTION	TARGET	PROPOSED BY
Executed	Roll back release rollback_release	subscription-app	ai-sre Review

AUDIT TRAIL

07:34:37.837 UTC	Incident opened detector	subscription-backend error rate 6% over the last minute, above its 3 sigma bound of 2.2% (mean 0.1% over 46 clean 10 s buckets)
07:38:01.169 UTC	Remediation proposed ai-sre	rollback_release subscription-app: Release v2 of subscription-backend (deployed on the subscription-app container) introduced a SELECT pg_sleep(0.8) query in the POST /api/subscribe endpoint. This slow query holds a database connection for ~800ms, exhausting the 20-connection pool under concurrent load, causing "connection pool exhausted: 20 of 20 connections in use (audit_trail)" errors on every new request. Traces 0a61bdd680e2489287c36aac23b120ad and 5cea99d2a7f3c8c424222a545e169807 confirm the pattern. act-2df618496285
07:38:33.521 UTC	Remediation refused ai-sre	Refused while pending: no human had approved it act-2df618496285

The audit trail: who proposed, who approved, what ran, what recovery was measured on, and when.

Room sees

- An incident naming the error rate and the bound it crossed.
- The agent naming release `v2` and `SELECT pg_sleep(0.8)` , citing trace ids, and proposing `rollback_release` to `v1` .
- The refusal in chat, and `remediation refused` in the timeline, before anyone approved.
- After Approve: the rollback command and its exit code, then `verification passed` and `incident resolved` , each with a timestamp and an actor.

Ask

Suppose the agent is right every time this week. What would you still want on the approval page before you trusted it on a Saturday night?

203 S

incident opened to proposal,
bad release

8

tool calls before it proposed

169 S

incident opened to proposal,
hung requests

\$0.0096

model cost for both investigations

Measured by `e2e/specs/12-report.spec.ts` on the last `make e2e` run with `deepseek/deepseek-v4-flash`, written to `evidence/e2e-report.json`. A test fails if these numbers drift from that file.

Concepts: the approval event, the audit log, the recovery check, autonomy levels.

17:00 to
17:30 ·
Module 6

What to harden before this runs on your systems

Goal Attendees leave with a list they can argue about with their own security team, and know which parts of it this lab already implements.

Say Everything below is a real gap or a real control in the lab in front of them. Go through it as an honest account of what is done and what is left, not as a checklist of advice.

Your telemetry carries personal data, and tool output goes to the model

The browser SDK is set to capture request bodies, so the email an attendee typed this morning is sitting in `SpanAttributes['http.request.body']`. Any tool that returns that span hands it to whoever hosts the model.

Today in the lab: the tools return trimmed rows, and the trace waterfall returns span attributes as they are. Nothing is redacted.

Redact at the tool boundary, not in the prompt. Return an allow-list of columns, hash identifiers so they still join, drop request and response bodies. Then point the same MCP tools at a model you host, which is one endpoint change in `librechat/librechat.yaml`, and decide per data class which model may see what.

Log lines are attacker-controlled text that your agent reads

Anyone who can get a string into your logs can write instructions to your agent. A user agent field is enough.

Today in the lab: the model can read any log body, and it can call only twelve tools. The one tool that acts takes an action and a target that must both be in a fixed allow-list, so an injected instruction has nothing to reach for.

Keep the action set closed and small. Treat every tool result as data, never as instructions, and say so in the system prompt. Alert on refusals: a proposal that fails grounding or policy is worth a human look, because it is either a confused model or a probe.

An approval page is only as good as what a tired person can check in a minute

At 3am the reviewer reads the root cause, which is the model's own prose, and decides. That is where a plausible wrong answer gets through.

Today in the lab: the trace ids, the service and the release are checked against ClickHouse before the proposal is stored, and the blast radius is our fixed text per action, not the model's. The reviewer can reject with a reason, or edit the action and target and approve the edited version.

Add what the reviewer cannot get from the page: a link that opens the cited traces in HyperDX, the diff of the release being withdrawn, who is on call for that service, and whether a change freeze is on. For anything above a blast radius you set, require a second approver.

The agent runs commands through the Docker socket

The control plane performs a rollback by running `docker compose up` for the previous image. That socket is the highest-risk surface in the lab.

Today in the lab: the command is built from a fixed list of two actions and never from model text, arguments are matched against a pattern, execution needs an approval event, and the VPS override publishes only Caddy's ports.

In production, do not give the agent a socket at all. Give it a narrow API that performs named changes, run that API as its own identity with its own audit, and let your existing deployment system do the deploy.

Runs cost money and can loop

Both investigations in the last run cost less than a cent, but an agent that retries a failing tool can spend without a ceiling.

Today in the lab: the prompt sets a budget of eight tool calls before proposing, and the report records tool calls, tokens and cost for every run.

Enforce the budget on the server, not in the prompt: cap tool calls per incident, cap spend per day, and stop the run when either is hit. Watch tool calls and time to proposal as a regression signal, the same way you watch latency.

An agent with no test set drifts silently

A prompt change, a model upgrade or a tool rename can all turn a working loop into a confident wrong one, and nothing in a chat window tells you.

Today in the lab: `make e2e` replays two incidents end to end against a real model and a real browser, and `make score` turns the result into a number out of 100 where every point is bound to a named test.

Keep your own incidents as fixtures. Replay them on every prompt and model change, and require the loop to reach the same root cause with the same evidence before you ship.

Ask

Which of these six would block this from running on your systems next quarter, and which would you accept with an audit log?

Concepts: goal hijack, context poisoning, trust exploitation, privilege abuse, evaluation sets.

Any time

If something goes wrong in front of the room

Fix

What you see	Why	What to do
No incident after two minutes of chaos	The baseline still holds a previous incident, or traffic is too thin	Run <code>make logs SERVICE=s1</code> talking; the 10% SLO ceiling c
The agent answers without calling tools	The conversation was opened fresh, so it has no incident id	Go back to the incident page than retyping the question.
A tool call fails with a schema error	LibreChat cached the tool definitions from before a rebuild	<code>docker compose -f docke</code> then reopen the conversation
The signup page hangs on Docs	Someone clicked the Docs link, which is the hang scenario	That is the Module 5 second i
HyperDX shows no rows	The time range is pinned to a window that has scrolled past	Type "Last 30 minutes" into th
Anything is stuck		<code>make chaos-reset</code> puts bot

17:30

Wrap-up

Say

The lab is gated the same way you would gate an agent in production. `make check` runs on every push. `make e2e` drives the loop you just watched with a real model and a real browser. `make score` turns both into a number out of 100, and every point is bound to a named test. Give them one instruction to take home: run the three substrate queries against their own observability data tomorrow. If those return nothing useful, no agent will either, and that is the first thing to fix.

Run

```
make chaos-reset  
make score  
make down
```

File: `scripts/rubric.py`

Concepts: the published reliability vocabulary, the named guardrails, the catalogued risks, what this lab chose for itself.