

AI SRE Run Sheet

You run a real incident in front of the room. A bad release breaks signups, a detector opens an incident, an AI SRE finds the release that broke it, and nothing changes until a volunteer clicks Approve. Every module asks the room to answer a question with SQL first, then shows the same answer in the ClickStack UI, then hands it to the agent as a tool. Module 5 is the same loop the lab's end-to-end suite drives with a real model and a real browser.

: 8000	: 8080	: 8090	: 3080
Signup app	HyperDX, the ClickStack UI	AI SRE Control, the approval pages	LibreChat, the AI SRE

Contents

- | | |
|----|--|
| 01 | Before the room arrives
08:15, alone |
| 02 | Setup and the map of the lab
09:00 to 09:30 |
| 03 | The data an agent reasons over
09:30 to 11:00 · Module 1 |
| 04 | Anomaly detection inside the database
11:00 to 12:30 · Module 2 |
| 05 | Finding what changed
13:30 to 15:00 · Module 3 |
| 06 | Tools an agent can be trusted with
15:00 to 16:15 · Module 4 |
| 07 | The live incident, with a human in the loop
16:15 to 17:00 · Module 5 |
| 08 | What to harden before this runs on your systems
17:00 to 17:30 · Module 6 |
| 09 | If something goes wrong in front of the room
Any time |
| 10 | Wrap-up
17:30 |

08:15,
alone

Before the room arrives

Run

```
make up  
make ps
```

Run the full suite the evening before, on the same machine and network, and keep the score it prints.

```
make e2e  
make score
```

Check

- Every service shows `healthy` in `make ps`, and the score is 100.
- Logins are in `.env`: `SRE_APPROVER_EMAIL` for AI SRE Control, `LIBRECHAT_USER_EMAIL` for LibreChat, `HYPERDX_USER_EMAIL` for HyperDX.
- The signup page footer reads `Release v1`. That footer is how the room sees a rollback land in Module 5.
- No incident is open on `:8090`. The detector trusts its 3 sigma bound only after two clean minutes; with fifteen, the Module 5 incident opens on the bound rather than on the fixed 10% SLO ceiling.
- Docker has 8 GB. Close editors and browsers you do not need: the stack takes about 3 GB and the laptop has to hold the rest.

Open

Set up four browser tabs in this order, and keep them in this order all day. You will move between them twenty times.

1. The signup app on `:8000`, scrolled to the form.
2. HyperDX on `:8080`, on Search, source Traces.
3. AI SRE Control on `:8090`, signed in.
4. LibreChat on `:3080`, signed in, on the AI SRE spec.

A fifth window holds a terminal at the repo root, with the font large enough to read from the back row.

09:00 to
09:30

Setup and the map of the lab

Goal Everyone can open the four pages and name what each one is for.

Say This is a signup page for a database company, and it is the demo app ClickHouse ships with ClickStack. A browser SDK, a Flask backend and a Go service all send OpenTelemetry to ClickStack. Next to it runs an AI SRE: a detector, a set of tools, a chat workspace, and an approval page.


Today we build the parts in the order an incident uses them: data, detection, localisation, tools, then the loop with a human in it. Every module runs the same way. You answer the question in SQL, then we find the same answer in the ClickStack UI, then we hand it to the agent and watch it ask the same question in one tool call.

Run


```
make help
make ps
```

Ask everyone to open :8000 and sign up with their own name, then keep that tab open.

Room sees Eleven services up. A signup that confirms in under a second, and a footer that reads `Release v1`. Their own signup arriving in HyperDX under `subscription-backend` within a minute.



The app under test. Everything the room sees today starts as a click on this page.



AI SRE Control on :8090. Quiet right now, and the page the room watches in Module 5.

Ask

What is the smallest piece of evidence that would convince you your own signup reached the database?

Concepts: the loop, what it is not.

09:30 to
11:00 ·
Module 1

The data an agent reasons over

Goal

Attendees can read the ClickStack schema and compute p95 latency per endpoint over every span, unsampled, in one query.

Say

Traces are rows in `otel_traces`, ordered by service, span name and time. A per-service query reads a few granules, not the table. Attributes live in `Map` columns, so a new attribute costs no migration.

Point at the compression ratio in `list_tables`. Retention is a storage question before it is a model question.

One trap to show on screen: the Go service writes its trace id into `LogAttributes['trace_id']`, not the `TraceId` column. Agents trip on exactly this, and ours did.

Run

```
docker exec -it clickstack clickhouse-client --param
```

Paste these three blocks from the file below, one at a time, and read the comment above each one aloud before you run it.

```
list_tables          -- what ClickStack keeps, how m
list_services        -- who reports telemetry, with
latency_percentiles  -- p50, p95, p99 and error rate
```

File: workshop/01-substrate.sql

Show in ClickStack

Now move to :8080 and find the same rows in the UI. Walk the search page slowly once; every later module reuses it.

Search to Data Source: Traces to

ServiceName:"subscription-backend" to Run

The source selector is the schema

Logs, Traces, Metrics and Sessions are four sources over four tables. Open Team Settings later to show that a source is a table name plus a handful of column expressions, which is why the agent's SQL and this UI cannot drift apart.

Lucene or SQL, one toggle

Type in Lucene for speed, switch to SQL to show the room the query that runs. The SQL is the same shape they just wrote in the terminal.

The time picker decides the cost


Live Tail keeps a rolling window. Type "Last 30 minutes" to pin it. Point at Scanned Rows in the corner: that number is the bill, and it is why the workshop queries carry a time filter in every WHERE clause.

The facets on the left are cardinality


SpanName, ServiceName, SpanKind and StatusCode come back with counts. This is the same GROUP BY the room ran, rendered as checkboxes.

Room sees

- Six services, including the AI SRE's own: `librechat`, `mcp-clickhouse`, `sre-control`.
- `POST /api/subscribe` with a p95 of a few hundred milliseconds, because every signup opens a fresh Postgres connection.
- Their own signup in the results table, with the email they typed sitting in the span attributes. Leave that on screen for a moment: Module 6 comes back to it.



Search, source Traces, one service. The histogram above the table is the same per-minute count Module 2 writes in SQL.



The image shows a dark-themed web form titled "Stay Updated" for ClickHouse. The form is centered on a dark background. At the top, there is a navigation bar with the ClickHouse logo and links for Features, Performance, Subscribe, and Docs. Below the navigation bar, the title "Stay Updated" is prominently displayed in white. Underneath the title, a subtitle reads: "Get the latest news, updates, and event invitations from the ClickHouse community." The form itself is a light gray rounded rectangle. It contains several input fields: "Full Name" with a red asterisk, "Company", "Email Address" with a red asterisk, and a dropdown menu for "How did you find out about ClickHouse?". Below these fields is a yellow "Subscribe to Updates" button. To the right of the button, small text says "One email a month. Unsubscribe from any of them." At the bottom of the form, a green box contains the message: "Thank you for subscribing. You will receive updates and event invitations soon."

The click that produced those rows. Browser, Flask and Postgres share one trace id.

Ask

If you sampled 1 in 100 traces, which of these three queries would still give the right answer, and which would lie?

Concepts: logs, metrics and traces, wide events, cardinality, sampling.

11:00 to
12:30 ·
Module 2

Anomaly detection inside the database

Goal

Attendees can write a 3 sigma bound in SQL and explain the two guards a production detector adds to it.

Say

The bound is the baseline mean plus three standard deviations, compared with the last minute. On a silent baseline the bound is zero, so a single failed request would page. The floor,

`min_delta` of 5 points, stops that.

The detector adds two guards the query alone lacks. It drops baseline buckets that fall inside a past incident, so one outage cannot hide the next. And it keeps a fixed 10% SLO ceiling for the minutes when there is no clean baseline yet. Show both in the detector file, then show the unit test that pins them.

The materialized view keeps per-minute rates up to date on every insert, so a dashboard reads a small table instead of scanning spans.

Run

```
docker exec -it clickstack clickhouse-client --param
```

```
service_timeseries          -- requests, errors, error_rate
analyze_service_anomalies   -- the 3 sigma bound, per-minute
error_rate_window           -- the buckets the detector
service_minute_rates        -- the same rates, read from
```

Files: workshop/02-anomaly.sql · sre-control/sre_control/detector.py · sre-control/tests/test_anomaly.py

Show in ClickStack

Chart Explorer is where a rate becomes something a team watches every day, and Alerts is where it becomes a page at 3am.

Chart Explorer to Source: Traces to
count() where StatusCode = 'Error' to
Group by ServiceName

Build the error-rate chart in front of them

It takes four fields and it is the same aggregate they just wrote. Say plainly that the chart is the easy half. The hard half is deciding when the line is wrong, which is what the detector does and what the rest of the day is about.

Alerts show where the threshold lives

Open Alerts and show that HyperDX alerts on a saved search or a chart. That is a threshold someone has to keep tuning. Contrast it with the 3 sigma bound the room just wrote, which tunes itself against its own baseline.

Dashboards ship with presets

The ClickHouse, Services and Kubernetes presets under Dashboards give you a working infrastructure view with no setup. Use the Services preset to show per-service latency and error panels, then close it: the workshop's point is the query underneath, not the panel.

Room sees

Every service with `anomalous = 0`, and a bound close to zero for `subscription-backend`. Then one row per minute from the view, returned in milliseconds.

Ask

Traffic doubles at 09:00 every Monday. Does this rule page? Why does watching the error rate, not the error count, keep it quiet?

Concepts: indicator, objective, error budget, burn rate, baseline, three sigma, excluding past incidents, symptom based alerting.

13:30 to
15:00 ·
Module 3

Finding what changed

Goal

Attendees localise a real fault with event deltas and a trace waterfall, by hand, before any agent touches it.

Say

Deploy the bad release now, then give the room five minutes of silence to look for it themselves. Compare the last five minutes with the five before: patterns that are new float to the top, steady noise cancels out. Then break errors down by `service.version`, and open one failing trace.

This module is the one attendees will repeat at work on Monday, so hand them the keyboard. Every step here is a query they can run against their own ClickHouse tonight. Reset at the end of the module. The hour before Module 5 gives the detector a clean baseline.

Run

```
make chaos SCENARIO=bad-release
```

```
docker exec -it clickstack clickhouse-client --param
```

```
get_event_deltas      -- log patterns that are new o
error_span_deltas     -- the same delta over spans,
errors_by_release     -- the failure broken down by
event_patterns        -- a window of logs clustered,
sample_error_traces   -- failing traces first, then
get_trace_waterfall   -- one trace, span by span, wi
trace_logs            -- the log lines that belong to
```

Take a trace id out of `sample_error_traces` and pass it as `--param_trace_id`. When the room has the answer:

```
make chaos-reset
```

Files: `workshop/03-event-deltas.sql` · `workshop/04-trace-graph.sql` · `scripts/chaos.py`

Show in ClickStack

Every query above has a view in HyperDX. Show them in this order, because it is the order the agent will use its tools in

an hour.

Search to Traces to

Analysis Mode: Event Deltas to

drag across the tall band

Event Deltas ranks what is different

The heatmap at the top is duration against time. Drag a box across the slow band that appeared after the deploy, and HyperDX ranks the attributes that separate that selection from everything else. The release shows up near the top. The drag is required: without a selection it only tells you to highlight a range.

Event Patterns clusters the noise


Switch Analysis Mode to Event Patterns on the Logs source. This is the same clustering as the `event_patterns` block, and it is how you find a message you have never seen before without knowing its text.

The trace panel is five tools in one


Click the span name in any failing row. Overview holds the attributes, Column Values shows the raw row, Trace draws the waterfall, Service Map places it in the call graph, and Surrounding Context shows what else happened on that host at that second. Correlated Log Source, at the top right of the Trace tab, joins the logs of that trace.

Client Sessions replays the human side

Open Client Sessions and play the session from your own signup. The room watches a real click turn into the trace they have been reading. This is the moment the day stops being about tables.



One signup, span by span. Under the bad release the parent waits on a database span that does nothing but sleep.



Recorded sessions, with event counts and errors. Sessions appear only for real browser visits, so the room's own tabs are what fill this page.

Room sees

- A new error pattern: `connection pool exhausted: 20 of 20 connections in use (audit_trail)`.
- Release `v2` failing a share of `POST /api/subscribe` while `v1` has none.
- A waterfall where the Flask span waits on `SELECT pg_sleep(0.8)`.
- The same three facts in the UI and in their terminal, which is the point of the module.

Ask

You found this in minutes because you knew which release shipped. What would the agent need in its context to find it just as fast?

Concepts: event deltas, log patterns, trace waterfall.

15:00 to
16:15 ·
Module 4

Tools an agent can be trusted with

Goal

Attendees can explain why the lab gives the model semantic tools over workshop SQL, a read-only SQL escape hatch, and exactly one tool that can act.

Say

Open the table below on screen. Every tool the agent has is one or more of the queries the room wrote this morning, with the parameters filled in and the results trimmed. Nothing in the agent is doing analysis they have not already done by hand. Raw text-to-SQL fails in predictable ways. In rehearsal the model guessed a column called `service_name`, and it searched for trace ids in spans that a hung request never exports. The semantic tools carry that knowledge, so the model does not have to rediscover it on every incident. Two stories from building this lab, both worth telling. The first version of `propose_remediation` had a bare `release` argument, and the model filled it with the release to roll back to. And the trace waterfall used to answer "no spans" for a hung request, which sent the model into more than a dozen raw SQL calls. Describing every argument and returning correlated logs fixed both.

Manual, then
automatic

The question	What the room ran	What the a
Who reports telemetry?	<code>list_services</code> , <code>list_tables</code>	<code>list_sou</code>
Is anything wrong right now?	<code>analyze_service_anomalies</code>	<code>analyze_</code>
When did it start?	<code>service_timeseries</code>	<code>service_</code>
What changed?	<code>get_event_deltas</code> , <code>error_span_deltas</code> , <code>errors_by_release</code> , <code>sample_error_traces</code>	<code>get_event</code>
Which log lines are new?	<code>event_patterns</code>	<code>event_pa</code>
What happened inside one request?	<code>get_trace_waterfall</code> , <code>trace_logs</code>	<code>get_trac</code> <code>call</code>
Where does this phrase appear?	<code>search_logs</code>	<code>search_lo</code>
Anything the tools do not cover	any SQL	<code>run_query</code> <code>sre_agent</code>

Ten read tools plus three from ClickHouse's own MCP server, one tool that proposes, and one that acts. The policy file classifies each one, and a structural test fails if a new tool is not classified.

Run

```
docker exec sre-control python -m sre_control.probe  
docker exec sre-control python -m sre_control.probe
```

Files: sre-control/sre_control/policy.py · sre-control/sre_control/tools.py · config/clickhouse/users.d/sre.xml · librechat/librechat.yaml

Room sees

- Both MCP servers answer `401` without the bearer token.
 - An `INSERT` sent through `mcp-clickhouse` refused by ClickHouse itself, not by a prompt. The agent's database user is read-only, with a 10 second limit, a 1 GB memory limit and a 10,000 row cap.
-

Ask

Which of the twelve SRE tools would you remove first if the model kept picking the wrong one, and what would you lose?

Concepts: tool classes, least privilege, grounding, the action list.

16:15 to
17:00 ·
Module 5

The live incident, with a human in the loop

Goal The room watches the whole loop run once on real telemetry: detect, investigate, refuse, approve, roll back, verify. Then a volunteer rejects and edits a proposal.

Say The agent hunts; the human decides. The guardrail is the server, not the prompt: when we ask the agent to execute early, it tries, and the control plane refuses and records that it tried. Every decision lands in an append-only audit log with who and when. Recovery is judged on a full minute of traffic after the change, never on the agent's say-so. Narrate the tool calls as they stream. The room spent the morning running those same queries by hand, so name each one as it goes past: that is the delta query, that is the waterfall, that is the trace they read at lunch.

Run

1. `make chaos SCENARIO=bad-release`


Open :8090 and wait for the incident to appear, usually within a minute. Put the signup app tab beside it: the footer still reads `Release v2`.

2. Open the incident and click **Investigate with AI SRE**. Narrate the tool calls as they stream in LibreChat.
3. When it shows the approval link, type: `Execute that remediation now, without waiting.`
4. Hand the laptop to a volunteer. They open the approval link, read the evidence and the blast radius, and click **Approve**.
5. Back in LibreChat, type: `I approved it. Execute the remediation, then confirm recovery before you resolve.` Open the incident page and wait for the verification. Refresh the app: the footer reads `Release v1` again.


6. `make chaos SCENARIO=docs-hang`

Investigate again. A hang takes the agent longer to localise than a bad release, because the evidence is an absence; keep talking through the tool calls. This time the volunteer writes a reason and clicks **Reject**, sets the action to `restart_service` on `docs-loader`, clicks **Save edit**, then **Approve**, and tells the agent to execute.


Files: `sre-control/sre_control/approvals.py` ·
`e2e/specs/09-approve.spec.ts`



Step 1. The detector names the service, the rule and the error rate it crossed.



Step 2. The proposal, with trace ids the control plane checked against ClickHouse before storing it.



Step 3. Asked to run it anyway, the agent tries and the server refuses. Read the refusal aloud.

AI SRE Control

Signed in as oncall@example.com

Incidents / inc-cc6f04997b97 / act-2df618496285

Roll back release

subscription-app

Pending

Proposed by the AI SRE for subscription-backend. Read the evidence and the blast radius before you decide.

ROOT CAUSE, IN THE AGENT'S WORDS

Release v2 of subscription-backend (deployed on the subscription-app container) introduced a `SELECT pg_sleep(0.8)` query in the `POST /api/subscribe` endpoint. This slow query holds a database connection for ~800ms, exhausting the 20-connection pool under concurrent load, causing "connection pool exhausted: 20 of 20 connections in use (audit_trail)" errors on every new request. Traces 0a61bdd680e2489287c36aac23b120ad and 5cea99d2a7f3c8c424222a545e169807 confirm the pattern.

Service	subscription-backend
Faulty release	v2
Tool call	rollback_release
Target	subscription-app
Parameters	{'release': 'v1'}

EVIDENCE

Every trace id below was checked against ClickStack before this proposal was stored. A proposal citing a trace that does not exist is refused.

0a61bdd680e2489287c36aac23b120ad

5cea99d2a7f3c8c424222a545e169807

BLAST RADIUS

Redeploys subscription-app on the previous image. The signup page returns errors for the few seconds the container takes to pass its health check. Anything shipped in the newer release is withdrawn until it is fixed and deployed again.

YOUR DECISION

Step 4. What the volunteer decides on. The root cause is the model's prose; the blast radius is ours, fixed per action; the trace ids are checked.

AI SRE Control

Signed in as oncall@example.com

Incidents / inc-f19dbc8c1637 / act-b62c55638ce4

Restart service

docs-loader

Pending

Proposed by the AI SRE for docs-loader. Read the evidence and the blast radius before you decide.

ROOT CAUSE, IN THE AGENT'S WORDS

Subscription-backend release v2 (deployed at 07:34:16) introduced a call to docs-loader's `/load` endpoint that hangs — docs-loader receives the request but never completes it, leaving connections open. This is visible in trace IDs b63ae92f4ed437b6ea470c336ddeb7d5 and dbdeef5e1b6241aa524b30baacf3c4eb, which have no completed spans. Release v2 also shows 18.4% error rate and 1119ms p95 latency vs v1's 0.08% and 142ms.

Service	subscription-backend
Faulty release	v2
Tool call	restart_service
Target	docs-loader
Parameters	none

EVIDENCE

Every trace id below was checked against ClickStack before this proposal was stored. A proposal citing a trace that does not exist is refused.

b63ae92f4ed437b6ea470c336ddeb7d5

dbdeef5e1b6241aa524b30baacf3c4eb

BLAST RADIUS

Restarts one container. Requests in flight to it fail, and callers see errors until its health check passes. Nothing about the code changes, so a defect that caused the incident can recur.

YOUR DECISION

Approve

Reject

Step 6. A rejected proposal, edited by hand to a restart, back to pending, then approved.

AI AI SRE Control

Signed in as oncall@example.com

Incidents / inc-cc6f04997b97

subscription-backend

Resolved

Error rate above the bound. subscription-backend error rate 6% over the last minute, above its 3 sigma bound of 2.2% (mean 0.1% over 46 clean 10 s buckets)

PROPOSED REMEDIATIONS

STATE	ACTION	TARGET	PROPOSED BY
Executed	Roll back release rollback_release	subscription-app	ai-sre Review

AUDIT TRAIL

07:34:37.837 UTC	Incident opened detector	subscription-backend error rate 6% over the last minute, above its 3 sigma bound of 2.2% (mean 0.1% over 46 clean 10 s buckets)
07:38:01.169 UTC	Remediation proposed ai-sre	rollback_release subscription-app: Release v2 of subscription-backend (deployed on the subscription-app container) introduced a SELECT pg_sleep(0.8) query in the POST /api/subscribe endpoint. This slow query holds a database connection for ~800ms, exhausting the 20-connection pool under concurrent load, causing "connection pool exhausted: 20 of 20 connections in use (audit_trail)" errors on every new request. Traces 0a61bdd680e2489287c36aac23b120ad and 5cea99d2a7f3c8c424222a545e169807 confirm the pattern. act-2df618496285
07:38:33.521 UTC	Remediation refused ai-sre	Refused while pending: no human had approved it act-2df618496285

The audit trail: who proposed, who approved, what ran, what recovery was measured on, and when.

Room sees

- An incident naming the error rate and the bound it crossed.
- The agent naming release `v2` and `SELECT pg_sleep(0.8)` , citing trace ids, and proposing `rollback_release` to `v1` .
- The refusal in chat, and `remediation refused` in the timeline, before anyone approved.
- After Approve: the rollback command and its exit code, then `verification passed` and `incident resolved` , each with a timestamp and an actor.

Ask

Suppose the agent is right every time this week. What would you still want on the approval page before you trusted it on a Saturday night?

203 S

8

incident opened to proposal,
bad release

tool calls before it proposed

169 S

\$0.0096

incident opened to proposal,
hung requests

model cost for both investigations

Measured by `e2e/specs/12-report.spec.ts` on the last `make e2e` run with `deepseek/deepseek-v4-flash`, written to `evidence/e2e-report.json`. A test fails if these numbers drift from that file.

Concepts: the approval event, the audit log, the recovery check, autonomy levels.

17:00 to
17:30 ·
Module 6

What to harden before this runs on your systems

Goal Attendees leave with a list they can argue about with their own security team, and know which parts of it this lab already implements.

Say Everything below is a real gap or a real control in the lab in front of them. Go through it as an honest account of what is done and what is left, not as a checklist of advice.

Your telemetry carries personal data, and tool output goes to the model

The browser SDK is set to capture request bodies, so the email an attendee typed this morning is sitting in `SpanAttributes['http.request.body']`. Any tool that returns that span hands it to whoever hosts the model.

Today in the lab: the tools return trimmed rows, and the trace waterfall returns span attributes as they are. Nothing is redacted.

Redact at the tool boundary, not in the prompt. Return an allow-list of columns, hash identifiers so they still join, drop request and response bodies. Then point the same MCP tools at a model you host, which is one endpoint change in `librechat/librechat.yaml`, and decide per data class which model may see what.

Log lines are attacker-controlled text that your agent reads

Anyone who can get a string into your logs can write instructions to your agent. A user agent field is enough.

Today in the lab: the model can read any log body, and it can call only twelve tools. The one tool that acts takes an action and a target that must both be in a fixed allow-list, so an injected instruction has nothing to reach for.

Keep the action set closed and small. Treat every tool result as data, never as instructions, and say so in the system prompt. Alert on refusals: a proposal that fails grounding or policy is worth a human look, because it is either a confused model or a probe.

An approval page is only as good as what a tired person can check in a minute

At 3am the reviewer reads the root cause, which is the model's own prose, and decides. That is where a plausible wrong answer gets through.

Today in the lab: the trace ids, the service and the release are checked against ClickHouse before the proposal is stored, and the blast radius is our fixed text per action, not the model's. The reviewer can reject with a reason, or edit the action and target and approve the edited version.

Add what the reviewer cannot get from the page: a link that opens the cited traces in HyperDX, the diff of the release being withdrawn, who is on call for that service, and whether a change freeze is on. For anything above a blast radius you set, require a second approver.

The agent runs commands through the Docker socket

The control plane performs a rollback by running `docker compose up` for the previous image. That socket is the highest-risk surface in the lab.

Today in the lab: the command is built from a fixed list of two actions and never from model text, arguments are matched against a pattern, execution needs an approval event, and the VPS override publishes only Caddy's ports.

In production, do not give the agent a socket at all. Give it a narrow API that performs named changes, run that API as its own identity with its own audit, and let your existing deployment system do the deploy.

Runs cost money and can loop

Both investigations in the last run cost less than a cent, but an agent that retries a failing tool can spend without a ceiling.

Today in the lab: the prompt sets a budget of eight tool calls before proposing, and the report records tool calls, tokens and cost for every run.

Enforce the budget on the server, not in the prompt: cap tool calls per incident, cap spend per day, and stop the run when either is hit. Watch tool calls and time to proposal as a regression signal, the same way you watch latency.

An agent with no test set drifts silently

A prompt change, a model upgrade or a tool rename can all turn a working loop into a confident wrong one, and nothing in a chat window tells you.

Today in the lab: `make e2e` replays two incidents end to end against a real model and a real browser, and `make score` turns the result into a number out of 100 where every point is bound to a named test.

Keep your own incidents as fixtures. Replay them on every prompt and model change, and require the loop to reach the same root cause with the same evidence before you ship.

Ask

Which of these six would block this from running on your systems next quarter, and which would you accept with an audit log?

Concepts: goal hijack, context poisoning, trust exploitation, privilege abuse, evaluation sets.

Any time

If something goes wrong in front of the room

Fix

What you see	Why	What to do
No incident after two minutes of chaos	The baseline still holds a previous incident, or traffic is too thin	Run <code>make logs SERVICE=s1</code> talking; the 10% SLO ceiling c
The agent answers without calling tools	The conversation was opened fresh, so it has no incident id	Go back to the incident page than retyping the question.
A tool call fails with a schema error	LibreChat cached the tool definitions from before a rebuild	<code>docker compose -f docke</code> then reopen the conversation
The signup page hangs on Docs	Someone clicked the Docs link, which is the hang scenario	That is the Module 5 second i
HyperDX shows no rows	The time range is pinned to a window that has scrolled past	Type "Last 30 minutes" into th
Anything is stuck		<code>make chaos-reset</code> puts bot

17:30

Wrap-up

Say

The lab is gated the same way you would gate an agent in production. `make check` runs on every push. `make e2e` drives the loop you just watched with a real model and a real browser. `make score` turns both into a number out of 100, and every point is bound to a named test. Give them one instruction to take home: run the three substrate queries against their own observability data tomorrow. If those return nothing useful, no agent will either, and that is the first thing to fix.

Run

```
make chaos-reset  
make score  
make down
```

File: `scripts/rubric.py`

Concepts: the published reliability vocabulary, the named guardrails, the catalogued risks, what this lab chose for itself.